

Preliminary results master thesis

Pablo Paramio Pérez

May 29, 2026

Abstract

The two MTU15 granularity reforms turn within-hour bidding into a strategic margin, and the dispatchable fleet exploits it. We see this directly in the bid curves: in critical hours (morning ramp and evening peak) where within-hour residual demand actually varies, ladders widen substantially — broadly under intraday reform (CCGT, hydro, pump-storage) and more selectively under the day-ahead reform (hydro, CCGT). Wind, whose output spans all hours rather than concentrating in peaks, shows little critical-flat differential.

Quantity-side, the cleared mix shifts in the direction the bid widening would predict: DA15 raises CCGT day-ahead auction-cleared volume substantially under joint 95% placebo-correction, and pump-storage redirects part of its daily generation toward those same critical hours. Price-side, ID15 cuts clearing prices substantially in both IDA and DA markets, both raw and net of a same-calendar 2024 placebo (conditional on wind, solar, and gas, so the drop survives controlling for renewable supply). The cross-market symmetry (both markets drop by the same amount on the same day, even though ID15 only changed *intraday* granularity) is the signature of rational anticipation: traders compress the day-ahead risk premium once they know they can re-trade in finer 15-min slots intraday.

1. Setup: outcomes, windows, and seasonal control

Takeaway. We measure how MTU15 reshapes three families of outcomes: per-curve bid shape; aggregate clearing price and cleared energy; and post-only within-hour dispersion. Identification compares *critical* clock-hours (those with the most within-hour load variation) to *flat* clock-hours (overnight; the within-day control) across each reform’s pre/post window. Outcomes whose pre/post comparison crosses different seasons — chiefly DA15 prices and cleared quantities — are additionally seasonally adjusted.

(Two) reforms. The Spanish wholesale electricity market shifted from a 60-minute to a 15-minute market time unit in two steps: **ID15** (2025-03-19, intraday auctions and continuous market) and **DA15** (2025-10-01, day-ahead). After each reform there are four periods per clock-hour; before, one. The *operación reforzada* regulatory stance (REE’s post-blackout reinforced operation, continuous from 2025-04-28) is a continuous confound that forces CCGT post-dayahead inclusion; tight pre/post windows around each reform date hold its level constant within each arm. Figure 1 summarises the full timeline, including other reforms that are not directly relevant for our research question but are important for context.




Figure 1: **Full reform-window timeline.** Five named windows partition the post-2024-06-14 sample; the pre-IDA window is the baseline regime. Dotted vertical lines mark the five reform dates that produce the window boundaries.

Variables. Each (unit u , date d , period p) sell offer is a tranche ladder $\{(p_k, q_k)\}_{k=1}^{K_{u,d,p}}$ where p_k is the bid price (EUR/MWh) and q_k the quantity in MW¹. Offers carry one of three structural types: *simple* (a bare price–quantity ladder), *minimum-income condition* (MIC), or *block* (all-or-nothing across periods). Most technologies use simple bids, but CCGT is the outlier with the bulk of its offers in MIC or block form (§ 4, § 5(d))². Let $\text{MCP}_{d,p}$ be the OMIE clearing price *of the market in which the offer was submitted* and h a bandwidth that isolates the bids in contention: tranches close enough to the clearing price to plausibly compete in the *competing tier* of the offer (CCGT’s two-tier structure makes this distinction explicit, § 7). We anchor h to the *contemporary contested margin*: h is computed as $p_{90} - p_{50}$ of the MCP distribution *in the analysis window of the comparison, market by market*. The resulting eight bandwidths (one per reform $\times$ {real, placebo} $\times$ {DA, IDA}) are summarised in the table below; bandwidth-robustness across wider h choices is reported in § 5(e).

Window	h in DA (EUR/MWh)	h in IDA (EUR/MWh)
ID15 real	50	62
ID15 placebo (2024)	45	46
DA15 real	50	58
DA15 placebo (2024)	45	49

Now define the in-band set, MW weights, and in-band weighted mean bid price:

$$\mathcal{B}_{u,d,p} = \{k : |p_k - \text{MCP}_{d,p}| \leq h\}, \quad w_k = \frac{q_k}{\sum_{j \in \mathcal{B}_{u,d,p}} q_j}, \quad \bar{p}_{u,d,p} = \sum_{k \in \mathcal{B}} w_k p_k.$$

Per-curve (in-band, MW-weighted; one value per (unit, date, period)):

$$\sigma_{p,u,d,p} = \sqrt{\sum_{k \in \mathcal{B}_{u,d,p}} w_k (p_k - \bar{p}_{u,d,p})^2}, \quad N_{\text{eff},u,d,p} = \left(\sum_{k \in \mathcal{B}_{u,d,p}} w_k^2 \right)^{-1}.$$

¹Following OMIE *Potencia* field in *ficheros OMIE v1.37* §5.1.4.2.

²Approximate sell-offer shares by structural type (DA15/ID15): wind, solar, hydro and pump-storage $\sim 100\%$ simple; nuclear 64% simple, 36% MIC; CCGT 18% simple, 50% MIC, 32% block.




Figure 2: **Bandwidth h and the in-band region.** The supply curve is the sorted sell-bid ladder for one quarter and one market. The market clears where supply meets demand, fixing the MCP. Bid-shape outcomes $(\sigma_p, N_{\text{eff}})$ are computed only on tranches in the in-band region $[\text{MCP} - h, \text{MCP} + h]$ — the competing tier where strategic action lives. The parked tier above the clearing-price ceiling is a non-strategic placeholder (§ 7) and is excluded by construction.

σ_p = price spread of the bid ladder within a single quarter; N_{eff} = effective tranche count of that quarter (1 = a single block, large = many equal-sized tranches).

Aggregate: $\text{MCP}_{d,p}$ is the OMIE clearing price (EUR/MWh); $Q_{d,p} = \sum_{u,k} q_{u,d,p,k}^{\text{cleared}}$ is the system DA cleared quantity per (date, period)³.

Within-hour dispersion (post-only, used to calibrate the critical/flat partition in § 2): let $m_q = \sum_{k \in \mathcal{B}_q} q_k$ and $\bar{m} = \frac{1}{4} \sum_q m_q$. For each (unit, date, clock-hour) we compute the standard deviation across the four quarters of the clock-hour for each of the five outcomes⁴:

$$\begin{aligned} D_{\text{price}} &= \text{sd}_q(\bar{p}_q), & D_{\text{qty}} &= \text{sd}_q(m_q)/\bar{m}, & \text{SD}_{\text{price}} &= \text{sd}_q(\text{MCP}_q), \\ D_{\sigma} &= \text{sd}_q(\sigma_{p,q}), & D_N &= \text{sd}_q(N_{\text{eff},q}). \end{aligned}$$

$D_{\text{price}}, D_{\text{qty}}, \text{SD}_{\text{price}}$ are across-quarter shifts in bid LEVEL, bid VOLUME, and system clearing prices; D_{σ}, D_N are across-quarter shifts in ladder SPREAD and TRANCHE COUNT. All five are zero under the MTU60-equivalent strategy (identical bids in every quarter), so any positive value post-MTU15 is a direct trace of within-hour discrimination.

³The OMIE assigned-power field is per-period instantaneous power (MW), invariant to period length, so one 15-min quarter and one 60-min hour at the same MW are directly comparable as rates. Aggregating MW across the four quarters of a clock-hour gives the stock of energy delivered in that clock-hour, $\text{MWh}_h = \sum_{q \in h} \text{MW}_q \Delta t_q$. Aggregate-quantity coefficients below are reported as MWh per clock-hour (numerically equal to the per-period mean MW because the integration is over 1 h, but the unit is energy, not power).

⁴We omit the u, h subscripts for readability.

Seasonality. Spanish electricity has strong intra-year seasonality. It bites on the level outcomes — clearing prices and per-tech cleared energy — which co-move with heating demand in winter, cooling demand in summer, and the hydrological cycle. Bid-shape outcomes $(\sigma_p, N_{\text{eff}})$ are MCP-centered and MW-weighted on the in-band ladder, so price-level seasonality does not transmit.

2. The critical/flat partition: where granular pricing has economic content

Takeaway. Within-hour discrimination has economic content only where within-hour residual demand actually varies. Critical hours (morning ramp + evening peak) are where MTU15 can bite; flat hours (overnight) are the within-day control. Post-MTU15 within-hour dispersion is much larger in critical than flat hours on every bid-level and bid-shape outcome we measure — the partition correctly identifies where the reform lives.

Intuition. Granular pricing only has economic content where within-hour residual demand⁵ actually varies; if a clock-hour is internally flat, the four quarters look the same and there is nothing to discriminate on. We measure within-hour variation by $\sigma_{\text{within}}(h, d) = \text{sd}_q[\text{load} - \text{wind} - \text{solar}]_{q,h,d}$ on 15-min residual-demand data⁶ and partition clock-hours by median σ_{within} :

- **Critical** $\mathcal{C} = \{5, \dots, 8, 16, \dots, 22\}$: morning ramp + evening peak; $\sigma_{\text{within}} \in [522, 1303]$ MW.
- **Flat** $\mathcal{F} = \{1, 2, 3\}$: overnight; $\sigma_{\text{within}} \in [336, 357]$ MW.
- **Midday** $\{11, \dots, 14\}$ (solar-marginal, $\sigma_{\text{within}} \in [468, 502]$ MW): excluded from main specs; kept for falsification (§ 5(a)).

Formal calibration of the partition (post-only). Let $D_{d,h}$ stand for any one of the five within-hour dispersion outcomes ($D_{\text{price}}, D_{\text{qty}}, \text{SD}_{\text{price}}, D_{\sigma}, D_N$) on date d , clock-hour h . Under MTU60 the within-hour quarter dimension does not exist, so $D_{d,h}$ is mechanically zero pre-reform; we use the post-MTU15 sample to *describe* where within-hour dispersion lives across hour-classes. A date-FE regression is a convenient way to summarise the critical–flat differential:

$$D_{d,h} = \delta_d + \theta \mathbf{1}\{h \in \mathcal{C}\} + \varepsilon,$$

fit per outcome. θ is the post-MTU15 critical–flat dispersion gap. This is a **descriptive contrast**, **not a causal estimate**.

⁵Spanish-market convention calls this the *hueco térmico* (thermal gap) — the load that thermal and other dispatchable generation must cover after subtracting wind and solar.

⁶ENTSO-E Transparency Platform: A65 (actual total load) and the wind/solar component of A75 (actual generation per type, PSR codes B16, B18, B19); window 2023-01-01 to 2026-05-14.

Outcome	Subsample	ID15	DA15	crit/flat ratio (IDA, DA)
D_{price} (EUR/MWh)	per unit (pooled)	0.39*** (0.05)	1.35*** (0.11)	4.1×, 7.4×
D_{qty} (CV)	per unit (pooled)	0.024*** (0.003)	0.033*** (0.002)	4.3×, 4.9×
SD_{price} (EUR/MWh)	system	4.55*** (0.40)	3.97*** (0.37)	2.8×, 2.6×
D_{σ} (EUR/MWh)	CCGT	−0.24 (0.35) ns	+0.57*** (0.10)	ns, 5.9×
	Hydro	+0.48*** (0.07)	+1.22*** (0.16)	2.2×, 3.1×
	Hydro pump	+0.16*** (0.03)	+0.66*** (0.15)	3.7×, 10.6×
	Wind	−0.003 (0.005) ns	+0.04*** (0.02)	1.0×, 2.3×

SE clustered by date; *** $p < 0.01$, ns = not significant. Bid-level rows (D_{price} , D_{qty} , SD_{price}) are bandwidth-free. The per-tech D_{σ} rows are computed at the new p_{90} baseline (§ 1): ID15 column uses post-ID15 IDA data with $h=62$; DA15 column uses post-DA15 DA data with $h=50$. *Reading.* Under DA15, D_{σ} is positive and significant for all three dispatchable techs in DA — CCGT +0.57, hydro +1.22, pump-storage +0.66 — with the largest critical/flat ratio for CCGT (5.9×). Under ID15, hydro and pump-storage widen their IDA within-hour ladders (+0.48 and +0.16), while CCGT IDA D_{σ} is not distinguishable from zero — consistent with combined-cycle gas mostly competing in DA rather than IDA at this stage. Wind stays small everywhere, the price-taking baseline.

3. Identification: DiD for bid shape, BSTS for prices and quantities

Takeaway. Bid shape is identified by a critical/flat DiD: the (critical – flat) shift across each cutover is the ATT under parallel trends. Clearing prices and per-tech cleared energy are identified by a Bayesian Structural Time Series counterfactual on a daily panel: conditional on the wind and solar renewable-supply covariates, the post-cutover series departs from the model’s pre-trained extrapolation by the treatment effect; a 2024 same-calendar placebo gauges residual cross-season trend.

u unit; d date; p period; h clock-hour; s IDA session; T_r reform date. Treatment status $D_{d,h} \equiv \mathbf{1}\{d \geq T_r\} \cdot \mathbf{1}\{h \in \mathcal{C}\}$; sample restricted to $h \in \mathcal{C} \cup \mathcal{F}$ for the DiD.

3.A. Spec A: Bayesian Structural Time Series for prices and per-tech cleared energy

Following Brodersen et al. (2015) as applied to the EPEX 15-min reform by Märkle-Hußet al. (2020), the daily response y_t (DA or IDA clearing price; per-tech daily cleared GWh) is modelled as a Gaussian state-space $y_t = \mu_t + \gamma_t + X_t' \beta + \varepsilon_t$, with μ_t a local-level random walk with stochastic slope, γ_t a 7-day weekly cycle, and $X_t = (\text{wind GWh}, \text{solar GWh}, \text{gas EUR})$ exogenous covariates under a spike-and-slab prior on β (gas added to Märkle-Hußet al.’s wind+solar; demand excluded to avoid simultaneity). Pre-cutover data train the model in **CausalImpact**; the counterfactual $\hat{y}_t(0)$ on the post-window is the extrapolation under realised X_t , and $\bar{\tau} = \frac{1}{|\text{post}|} \sum_{t \geq T_r} (y_t - \hat{y}_t(0))$ with its 95% credible interval is read from the MCMC draws. Cleared-quantity outcomes are reported in absolute GWh/day rather than % (the Gaussian state-space can produce negative posterior draws when covariates predict very low CCGT dispatch, and %-readings would re-introduce log-like unit-dependence issues, Chen and Roth 2024).

Windows hold the relevant regulatory regime constant across pre and post. DA15: pre 2025-04-28 → 2025-09-30 (156 d, starting at reforzada onset); post 2025-10-01 → 2025-11-09. ID15: pre 2024-06-14 → 2025-03-18 (278 d, post-IDA-reform 3-session architecture); post 2025-03-19 → 2025-04-27 (40 d, pre-blackout). A 2024 same-calendar placebo for each reform falsifies residual cross-season trend; placebo-net effects are the cleanest identification of the treatment. *Identifying restriction*: conditional on X_t , $y_t(0)$ post-cutover equals the pre-trained extrapolation; the placebo gauges whether that extrapolation is itself biased by seasonality.

3.B. Spec B: per-hour-class BSTS on bid-level outcomes

Spec C below identifies the within-quarter shift in bid *shape*. To identify within-day reallocation in bid *levels*, we run Spec A-style BSTS *independently* on each (tech, market, hour-class) cell for two bid-level outcomes: $\bar{p}_{t,m,h,d}$ (MW-weighted mean of in-band tranche prices) and $m_{q;t,m,h,d}$ (total in-band quantity, summed across tranches/periods/units and for IDA also sessions). Hour-classes follow the calibration of § 2; same wind+solar+gas covariates and same windows as § 3.A; the in-band bandwidth is window-and-market-specific $h = p_{90} - p_{50}$ of MCP (eight values in § 1). Within each (tech, market, outcome) the identified target is the (critical – flat) placebo-net contrast, post-processed from four BSTS posteriors (real/placebo × critical/flat); the joint 95% credible interval treats the four posteriors as independent. We run BSTS per hour-class and form the contrast in post-processing (rather than apply BSTS to the differenced series) so the seasonal structure stays available for the model to absorb.

3.C. Spec C: per-curve DiD for bid shape

For $Y \in \{\sigma_p, N_{\text{eff}}\}$, per tech:

$$Y_{u,d,p[s]} = \alpha_u + \beta \mathbf{1}\{d \geq T_r\} + \xi \mathbf{1}\{h \in \mathcal{C}\} + \theta D_{d,h} + \varepsilon, \quad \text{SEs clustered by date.}$$

We use DiD rather than BSTS here because the bid-shape outcomes are MCP-centered and MW-weighted (§ 1), so they do not inherit price-level seasonality. Differencing absorbs any common-hour seasonality, and under parallel trends on the (critical – flat) differential, θ identifies the ATT on critical-hour bid shape. Flat hours act as a low-exposure within-day control (mechanically reachable post-cutover but with negligible within-hour residual demand variance, § 2), giving the standard 2×2 DiD without staggered-treatment negative weights (de Chaisemartin and D’Haultfoe uille, 2026, ch. 3).

4. Results: bid-shape widening, day-ahead CCGT auction-cleared scale-up, and directional price movements

We report findings one spec at a time, A then B then C, mirroring the order of § 3: [Spec A](#) (Bayesian counterfactual on daily prices and per-tech cleared energy, § 3.A), [Spec B](#) (per-hour-class BSTS on bid-level outcomes — mean in-band price and total in-band quantity, § 3.B), and [Spec C](#) (per-curve DiD on bid shape, § 3.C). Spec A identifies daily-level effects on clearing prices and per-tech cleared energy; Spec B identifies within-day reallocation of bid *levels* across hour-classes; Spec C identifies within-quarter bid-*shape* effects. All three share the same wind+solar+gas covariates and the 2024

same-calendar placebo. An integrative takeaway closes the section.

4.A. Spec A: prices and per-tech cleared energy

Outcome	Tech	ID15		DA15	
		Effect	Placebo	Effect	Placebo
Clearing price (EUR/MWh)	—	−38.34***	−18.08*	+2.97	−8.54
Cleared (GWh/day)	CCGT	−8.08***	−9.18***	+40.31***	−18.69***
	Hydro	−1.27***	+4.98***	+6.16*	+9.28***
	Hydro pump	−0.37*	+0.54**	+4.43***	+1.42*
	Wind	+5.34***	+7.99***	+13.16***	+2.15
	Solar	+6.70***	+2.55***	−3.37	+1.38
	Nuclear	+11.51***	+19.06***	−18.21**	−34.77***

Bayesian posterior point estimates; tail-probability stars * < 0.10, ** < 0.05, *** < 0.01. Windows in § 3.A; placebo columns are the same calendar window in 2024 with a fake cutover. The placebo-net effect (Effect − Placebo) is the identified treatment; the joint 95% credible interval propagates both posterior variances. Observed-vs-counterfactual series in Figures 3–5.




Figure 3: **ID15 effects on both markets** (2×2 zoomed view, 40-day pre-blackout post-window). Black: observed (7-day MA). Blue: BSTS in-sample fit on the pre-window with 95% band. Red dashed: out-of-sample counterfactual with 95% band. Cutover 2025-03-19; blackout 2025-04-28. Both IDA and DA clearing prices (top row) drop below the counterfactual by similar magnitudes on the same day, even though ID15 only changed intraday granularity — the cross-market rational-anticipation channel into day-ahead (finding (ii)).




Figure 4: **DA15 effects on both markets** (2×2 zoomed view, post-window through 2026-02-26). Same colour scheme as Figure 3. Cutover 2025-10-01; reforzada-constant pre-window starts 2025-04-28. The DA CCGT panel (bottom-right) is the one quantity outcome where the observed series visibly departs from the counterfactual post-cutover — the surviving CCGT scale-up of finding (iii). By contrast, the IDA panels show no backwards transmission (intraday clearing follows day-ahead, so a DA-side reform cannot propagate backwards).

(iii) **[Surviving] Day-ahead combined-cycle gas auction-cleared volume scales up under DA15.** The headline magnitude is in the BSTS table; this is the *only* level-shift finding that survives the conservative joint 95% test net of the 2024 same-calendar placebo. The channel is the same as the bid-shape widening in (i): MTU15-DA lets combined-cycle gas discriminate finely within the hour, raising its expected acceptance probability across the four quarters of a peak hour. Extending the post-window through the full available winter shrinks the estimate (Figure 4, bottom-right) — still positive, but pulled down as observed cleared volume mean-reverts after a strong October. Hydro, solar and nuclear are essentially null net of the 2024 placebo. **[NOVEL] The CCGT cleared scale-up is one-sided: substantial on the day-ahead side, essentially unchanged on the intraday side (per-tech cross-market BSTS). Combined-cycle gas migrates into DA without losing the IDA arm in cleared terms, so the bid-side ladder migration of (v) does not translate symmetrically to clearing.**

(ii) **[Wounded] ID15 raw clearing-price drop in both markets, but the conservative joint test does not retain.** The raw drop is the cleanest reform-attributable signal in the price data (BSTS table; top row of Figure 3). Three credibility signals stack: the BSTS already conditions on wind, solar and gas, so the drop is what remains after the merit-order shift; both markets drop by the same magnitude on the same day even though ID15 only changed *intraday* granularity, the rational-anticipation channel into day-ahead; and across the pre-window the BSTS hugs the observed series tightly (Figure 3). The placebo-net joint 95% credible intervals do bracket zero, but this is a conservative test that propagates two posterior variances; the placebo’s own loose CI reflects recurring spring-month model residual rather than a clean null. The DA15 price effect, by contrast, is small and not directional. *Verdict.* The finding is wounded by the 2024 same-calendar placebo (§ 5(c)): the placebo-net joint test does not retain it, so we report the raw drop and its rational-anticipation signature without claiming a surviving level effect. **[NOVEL] Wedge volatility under ID15**

jumps in critical hours. The wedge level (BSTS on the daily DA – IDA wedge) is null in both reforms. The within-day SD of the wedge across the 24 clock-hours, by contrast, jumps sharply under ID15 and concentrates in critical hours, robust across both a 2024 and an independent 2026 same-calendar placebo. The intuition is direct: ID15 makes intraday 15-min while day-ahead stays 60-min, so each clock-hour hosts four IDA quarter-prices against one DA price, and the per-hour wedge moves more where within-hour residual demand actually varies. The symmetric reform (DA15, both markets at 15-min) does *not* revert the wedge SD — residual candidates are behavioural hysteresis or a renewable-penetration trend the covariates do not fully absorb. ISP15 (settlement-period reform, OMIE contract grid unchanged) points the same way as ID15 — consistent with bidders responding to a finer settlement reference even without a grid change.⁷

(iv) Pump-storage daily cleared MW: ID15 intraday-auction decrease surviving; DA15 day-ahead rise borderline. Figure 5 shows both effects against their 2024 placebos. The tech with the largest within-day arbitrage value reallocates its dispatch once each market becomes 15-min-priceable. ID15 pulls pump-storage capacity *out of* the now-finer intraday auction (top-left below the band, top-right placebo flat; placebo-net survives joint 95%). DA15 lets the day-ahead auction absorb more of its discharge (bottom-left above the band) — the closed-form storage prediction (§ 3) — but the placebo (bottom-right) also rises, so the daily-level test is borderline; the within-day redistribution shows up cleanly in the bid-shape geometry of (i) and the per-hour-class quantity migration of (v) instead.




Figure 5: **BSTS counterfactual for pump-storage cleared MW.** *Top row:* ID15 effect on intraday-auction pump-storage cleared (left) and the 2024 same-calendar placebo (right). *Bottom row:* DA15 effect on day-ahead pump-storage cleared (left) and its 2024 placebo (right). Colour scheme as Figure ??.

⁷Wedge SD placebo-net under ID15 is +2.16 EUR/MWh (joint 95% [+0.62, +3.70]) overall and +3.14 in critical hours ($z \approx 4.4$). DA15 placebo-net is null. ISP15 placebo-net is +3.92.

4.B. Spec B: within-day reallocation under MTU15 (per-hour-class BSTS)

Figure 6 reports the *placebo-net level* of the BSTS effect for each hour-class (critical, midday, flat) separately, rather than just the (critical – flat) differential. The differential collapses information — two hour-classes can move in the same direction at different magnitudes, and the differential then hides the level story. The level view also makes explicit a caveat: these outcomes are *bid-level* ($\bar{p}$ and $\sum m_q$ are computed on submitted bids in the window-and-market-specific $\pm h$ EUR/MWh band around MCP, with $h \in \{45, 46, 49, 50, 58, 62\}$ as set in § 1), not cleared/dispatched quantities. Cleared quantity is captured by [Spec A](#) in § 4.A; the two can move in opposite directions (e.g., under ID15 CCGT’s BSTS daily cleared MW falls while its in-band offered MW rises, because more tranches qualify as “in-band” when MCP itself falls).




Figure 6: **Per-hour-class BSTS on bid-level outcomes — placebo-net levels.** Each row is one BSTS run. The (tech, market) group is bold-labelled on the right; within each group, three rows give the placebo-net effect for critical (red), midday (orange) and flat (blue) hours separately. Error bars are the 95% credible interval propagating real and placebo posterior variances. *Faded markers and bars indicate the 95% CI brackets zero (effect not statistically distinguishable from zero).* *Top row of panels:* mean in-band bid price $\bar{p}$ (EUR/MWh). *Bottom row of panels:* total in-band quantity $\sum m_q$ (thousands of MW, summed across tranches, periods, units, and — for IDA — sessions). *Left column:* ID15. *Right column:* DA15. *Caveat:* these are *submitted-bid* outcomes, not cleared/dispatched MW; the in-band test (window-and-market-specific $\pm h$, $h \in \{45, 46, 49, 50, 58, 62\}$; see § 1) is MCP-relative, so a drop in MCP can mechanically pull more tranches inside the band without any change in bidding behaviour.

(v) [Surviving] Under ID15 the dispatchable fleet migrates in-band offering *into* IDA across hour-classes and bid prices fall; under DA15 the same fleet migrates from IDA *into* DA across hour-classes and bid prices rise. (Magnitudes in Figure 6; surviving cells are full-colour markers.) The substitution picture is sharp on both outcomes:

- *In-band quantity*: under ID15, in-band volume grows in IDA across all three hour-classes for combined-cycle gas, hydro and pump-storage; under DA15, in-band volume grows in DA across hour-classes and falls in IDA. The shift is comprehensive across the day, not concentrated in critical hours — dispatchables migrate the *entire competing-tier ladder* into the market that has just become granular.
- *Bid-price level*: ID15 drives placebo-net mean in-band bid prices substantially down across cells (sharper in flat/midday); DA15 raises them across cells. The placebo-net contrast strips the calendar trend; the surviving effects identify firm-side restructuring of competing-tier bid levels.

Bid vs. clearing caveat. Granularity in market m attracts in-band *bidding* into m (the upstream-migration channel the model in § 3 predicts). Whether the migration translates to cleared MW is a separate question answered by Spec A: DA15 CCGT cleared MW does scale up (finding (iii)); ID15 CCGT cleared MW *drops* even as in-band offered MW expands — the same MCP-relative artefact that lets more tranches fall “in-band” when MCP itself falls.

Tech selection for Spec C. CCGT (the focal scarcity-tier technology), hydro and hydro pump for the dispatchable comparison, and wind as a price-taking falsifier.⁸ Solar PV is excluded (zero in-band scarcity-tier capacity); nuclear is excluded (mostly bilateral). Hydro, hydro pump and wind are $\sim 100\%$ simple offers in both arms, so their Spec C DiDs are simple-only by construction; the offer-type composition concern is CCGT-specific.

4.C. Spec C: bid-shape widening in critical hours

Outcome (DiD θ)	Subsample	ID15		DA15	
		DA ($h=50$)	IDA ($h=62$)	DA ($h=50$)	IDA ($h=58$)
σ_p (EUR/MWh)	CCGT	+4.17*** (0.44)	+0.10 (0.53)	+0.68*** (0.22)	−0.68* (0.36)
	Hydro	−0.53* (0.29)	+0.43* (0.25)	+1.80*** (0.34)	+0.30 (0.25)
	Hydro pump	−0.73*** (0.27)	−0.01 (0.18)	+0.92*** (0.35)	+0.32* (0.19)
	Wind [†]	−0.10*** (0.03)	+0.04** (0.02)	−0.02 (0.03)	−0.02*** (0.003)
N_{eff}	CCGT	−0.08 (0.21)	+0.79 (0.73)	+1.89*** (0.15)	−0.17 (0.53)
	Hydro	−0.24*** (0.06)	+0.18*** (0.05)	+0.51*** (0.09)	+0.00 (0.06)
	Hydro pump	+0.07 (0.05)	+0.12* (0.07)	−0.06 (0.07)	−0.16 (0.13)
	Wind [†]	+0.01 (0.02)	+0.03*** (0.01)	−0.01 (0.01)	−0.01 (0.01)

Notes: SEs in parentheses, clustered by date. $*p < 0.10$, $**p < 0.05$, $***p < 0.01$. Per-(tech, market) subsample regressions, unit FE absorbed via within-unit demeaning. Bandwidth h is window-and-market-specific (the p90 of the

⁸Wind has the dominant sell-side quantity but is not the marginal price-setter; CNMC (2025c) defines the marginal technologies as coal + CCGT + hydro + 5% RECORE. Aggregator portfolios construct dense single-price blocks at low EUR/MWh, so the clearing price lands inside a block and EUPHEMIA partially accepts a wind unit as the mechanical residual — not strategic price-setting. A strategically inert technology should yield small DiDs, validating the contrast.

MCP distribution in that window, market by market): see § 1 and § 3.C. Sensitivity to wider bandwidths is reported in § 5(e). [†]Wind reads as a price-taking baseline: 84% of in-band wind sell curves are a single block; magnitudes are an order of magnitude smaller than the dispatchable techs.

(i) Bid-shape widening in critical hours is market-specific: ID15 widens day-ahead CCGT and intraday N_{eff} ; DA15 widens day-ahead σ_p/N_{eff} broadly and tightens IDA CCGT. (Magnitudes in the table; Figure 7.) Under ID15, the granularity instrument lands on intraday auctions: N_{eff} widens significantly in IDA for hydro, pump-storage, and wind (more granular ladders), while σ_p on IDA stays small for combined-cycle gas. The cross-market spillover lands on day-ahead: CCGT widens its DA σ_p substantially (the anticipation-channel restructuring of DA ladders ahead of finer IDA re-trading), and hydro and pump-storage tighten their DA σ_p instead. Under DA15, the granularity instrument lands on day-ahead: σ_p and N_{eff} widen across all three dispatchable techs in DA, while CCGT IDA σ_p *tightens* — combined-cycle gas pulling its structured ladder out of IDA into DA in critical hours, consistent with the Spec B levels in (v). Wind reads as a price-taking baseline throughout.

Behaviour change, by reform.

ID15 (intraday-auction granularity, 2025-03-19). The dispatchable fleet pushes its bid ladder *into* IDA, where the granularity instrument now lives. In-band IDA quantity grows across all three hour-classes for CCGT, hydro and pump-storage ((v)); IDA N_{eff} widens for hydro, pump-storage and wind ((i)); the cross-market spillover lands on day-ahead CCGT σ_p , restructuring the DA ladder ahead of finer IDA re-trading ((i)). Pump-storage *leaves* the now-granular IDA channel in cleared-volume terms ((iv), surviving). Both clearing prices drop on the cutover and remain below the counterfactual for the 40-day pre-blackout window, but the joint placebo-net test does not retain the level effect ((ii)).

DA15 (day-ahead granularity, 2025-10-01). The same fleet pulls its bid ladder *out of* IDA *into* DA, mirroring the ID15 pattern. In-band DA quantity grows across hour-classes for all three dispatchables; in-band IDA quantity falls ((v)). DA σ_p widens for all three dispatchables and CCGT N_{eff} widens broadly; CCGT IDA σ_p *tightens* as combined-cycle gas pulls its structured ladder out of IDA into DA ((i)). Day-ahead CCGT cleared volume scales up substantially (the only surviving level effect, (iii)). Day-ahead pump-storage cleared rises directionally but is borderline at joint 95% ((iv)).

Common mechanism. Both reforms produce the same upstream-migration of the dispatchable fleet's bidding into the market that has just become 15-min-priceable — the closed-form sequential-market model in § 3 predicts the signs and the per-tech comparative statics.




Figure 7: **Spec C per-curve DiD on bid-shape outcomes at the window-specific p90 bandwidth.** *Top row:* σ_p DiD (within-quarter price SD); *Bottom row:* N_{eff} DiD (effective tranche count). *Left column:* ID15. *Right column:* DA15. Each row inside a panel is one (tech, market) cell; blue = day-ahead, red = intraday auctions. Faded markers indicate the 95% CI brackets zero. Bandwidth h varies by (window, market): ID15 uses $h=50$ (DA) / $h=62$ (IDA); DA15 uses $h=50$ (DA) / $h=58$ (IDA). The picture is genuinely market-specific: ID15's effects concentrate in IDA N_{eff} (more granular intraday ladders) and DA CCGT σ_p (anticipation-channel widening); DA15's effects concentrate in DA across all dispatchables, with CCGT IDA σ_p tightening (substitution into DA).

5. Robustness: surviving findings hold; price effects do not

Takeaway. Bid-shape findings survive the standard threats; the BSTS DA15 combined-cycle gas cleared scale-up survives at joint 95%; price effects do not. The single fragile bid-shape outcome at the p_{90} baseline is DA15 day-ahead pump-storage σ_p (pre-trend placebo exceeds the headline DiD; see § 5(b)); CCGT σ_p and N_{eff} are robust.

Pre-trend placebo — the headline check. A pre-only midpoint placebo splits each reform's pre-window at its midpoint and runs a fake critical-flat DiD on the same outcome. The intuition: if a placebo on the pre-window already finds a critical-flat shift, then what the real DiD picks up may be extrapolation of a pre-existing trend rather than a reform effect. ID15 IDA dispatchable σ_p placebos all bracket zero — the ID15 bid-shape findings rest on a clean pre-trend. For DA15 day-ahead: hydro

σ_p placebo is clean, CCGT σ_p placebo runs opposite the headline (so de-trending strengthens the real result), and pump-storage σ_p is the one wounded outcome — its pre-trend placebo exceeds the headline DiD, so we read it as not separable from a pre-existing trend.

Hour-class robustness. A midday-vs-flat falsification gives sign-stable but smaller-magnitude DiDs at the p_{90} baseline: the reform reshapes the entire non-flat zone, not just critical hours — consistent with the per-hour-class BSTS picture (§ 4.B). The cleaner critical-vs-flat contrast remains the headline because flat hours are the unambiguous within-day control (no within-hour residual-demand variation to discriminate on).

Cross-border and reforzada — the two natural confounds. The cross-border-timing channel (SIDC/XBID also moved to 15-min on the same date as ID15) is a sibling treatment on the same reform, not an omitted confound: adding hourly cross-border net flow as a per-curve control on ID15 IDA σ_p leaves the per-tech DiDs essentially unchanged. The reforzada regulatory stance, present in both arms of DA15, is also not an OVB: it acts on the *level* of Fase I up-redispatch (substantial jump in flat overnight hours) but not on its *within-hour SD*. The interpretation is mechanical: reforzada parks capacity in the scarcity tier and recalls it through restriction offers (§ 7), while MTU15 reshapes the competing tier within the band — two non-overlapping mechanisms.

Per-CCGT-firm decomposition of the DA15 σ_p widening. The pooled DA15 CCGT day-ahead σ_p widening is driven primarily by Naturgy and to a lesser extent Iberdrola; Endesa and EDP-HC do not contribute. The Naturgy effect is not a flat-baseline artefact at the p_{90} bandwidth: σ_p rises in both arms but more in critical hours.

Bandwidth h . Why p_{90} : it isolates the competing tier of CCGT bids and excludes the parked tier above the ceiling, which is a non-strategic placeholder rather than a strategic price (§ 7). Including the parked tier would bias σ_p upward by a measurement-driven cross-tier spread instead of a behavioural signal. The qualitative picture is robust to small shifts around p_{90} . At much wider bandwidths, CCGT σ_p magnitudes inflate substantially and several DiD coefficients flip sign because the cross-tier spread itself moves across the cutover for reasons unrelated to within-tier bidding — exactly the artefact p_{90} avoids.

Pump-storage seasonal control (Spec A). The DA15 day-ahead pump-storage rise is borderline at joint 95% in the main BSTS because of cross-season baseline drift between summer pre and autumn post. A same-calendar DiD restricted to Oct–Dec 2024 vs. Oct–Dec 2025 and a Fourier deseasonalisation both retain the directional finding when the demand-mix story is held fixed.

6. Continuous-market response at the three structural reforms

Takeaway. The continuous (XBID) market’s structural response is concentrated at ID15: trade count surges at the cutover — the mechanical effect of the contract clock-grid moving from 60-min to 15-min in delivery. Total energy turnover *falls* (smaller average trade), and the volume-weighted mean trade price drops sharply, mirroring the auction-clearing-price drop in § 4(ii). DA15 produces a smaller secondary effect (daily GWh up, borderline trade-count rise, no price effect). ISP15 (the December 2024 settlement-period reform) leaves continuous-market volumes and prices unchanged — consistent with its scope being settlement, not contract structure.

Specification. BSTS-style counterfactual on daily ES-leg continuous-market series (one of `seller_zone` or `buyer_zone` = 10YES-REE-----0), with wind+solar+gas covariates and a weekly seasonal cycle (same setup as § 3.A). *Units.* All quantity outcomes are *energy* (MWh = MW × delivery hours = `quantity_mw` × `mtu_minutes`/60, summed per day), not raw MW sums: cross-reform comparability requires the stock (energy delivered), not the rate, since post-MTU15 each clock-hour holds four 15-min contracts where pre-MTU15 it held one 60-min contract. The volume-weighted mean price is energy-weighted (weight = MWh of each trade). Three reform cutovers: *ISP15* (settlement period 60 → 15 min, 2024-12-11; pre 2024-06-14 → 2024-12-10 holds the IDA reform constant, post 2024-12-11 → 2025-03-18 stops at ID15); *ID15* (intraday auctions + continuous market → MTU15, 2025-03-19; pre 2024-12-11 → 2025-03-18 post-ISP15, post 2025-03-19 → 2025-04-27 pre-blackout); *DA15* (day-ahead → MTU15, 2025-10-01; pre 2025-04-28 → 2025-09-30 post-blackout, post 2025-10-01 → 2025-12-31). Each real arm has a 2023/2024 same-calendar placebo with a fake cutover.

Reform	Outcome	Real	Placebo	Placebo-net
ISP15	$n_{\text{trades}}/\text{day}$	−3,120***	−1,231***	−1,889 (brackets 0)
	GWh/day	−3.85*	−2.23**	−1.62 (brackets 0)
	VW-price (EUR/MWh)	−15.1 ns	−27.4***	+12.3 (brackets 0)
ID15	$n_{\text{trades}}/\text{day}$	+22,335***	+1,381*	+20,954***
	GWh/day	−7.19***	+6.72***	−13.91***
	VW-price (EUR/MWh)	−80.1***	−14.1***	−66.0***
DA15	$n_{\text{trades}}/\text{day}$	+5,018**	+257 ns	+4,761 (borderline)
	GWh/day	+3.85***	−1.35 ns	+5.20**
	VW-price (EUR/MWh)	+4.7 ns	−4.7 ns	+9.4 ns

Notes: BSTS posterior means; tail-probability stars * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$, ns = posterior interval brackets zero. Placebo-net joint significance based on independent posteriors propagating two variances.

(vi) ISP15: settlement scope, no contract-structure response. The December 2024 reform changed how REE settles imbalance positions (60→15-min ISP), but it did not change the OMIE continuous-market contract grid: traders still bought and sold hourly delivery contracts before and after. A null effect is therefore the correct prediction. Trade count drops 22% raw and energy turnover drops 14% raw, but the same drift is in the 2023 placebo (settlement bookkeeping innovations are not the source) and all three placebo-net intervals bracket zero. We report this as a methodological

sanity check: a settlement-only reform shows nothing on contract-structure outcomes, consistent with our specification’s reading of where the reform lives.

(vii) ID15: the contract grid moves from 60 to 15 minutes — the structural reform.

Pre-reform, each clock-hour of delivery hosted a single 60-min continuous-market contract; post-reform, four separate 15-min contracts. Three of our four outcomes pop accordingly:

- *Trade count.* Each hour now offers four matching opportunities where there used to be one, so the trade-count surge is mechanical at first order (placebo-net well clear of zero). Even “mechanical” is informative — it confirms market participants *used* the new finer grid rather than continuing to trade only on the hourly aggregate.
- *Energy turnover.* The placebo-net is *negative*: more contracts, but each smaller, and the total energy changing hands drops. Behavioural read: the granularity reform did not pull additional energy into the continuous market; it sliced the same (in fact slightly less) energy into smaller pieces. The drop is partly because some hourly counterparties drop out when they don’t see a single 60-min match on offer, and partly because dispatchables migrated bidding into the now-finer IDA auctions instead (§ 4(v)).
- *VW-price.* The continuous market clears in equilibrium with the auctions, and we see the same ID15 day-of-cutover clearing-price drop here as on auction-cleared prices, well clear of zero — the cross-market spillover documented for ID15 in § 4(ii).

(viii) DA15: a smaller, indirect continuous-market effect. Day-ahead granularity does not change the continuous-market contract grid (already 15-min post-ID15); only the day-ahead *reference price* that continuous traders watch becomes finer. With a finer DA reference, traders reposition more aggressively in continuous: daily GWh up (joint 95% retains), trade count borderline. The price effect is statistical noise. The interpretation is that DA15’s continuous-market footprint is incremental and proportional to how much tighter the day-ahead reference now is, not a structural break — exactly as expected from a reform that touches the reference market but not the continuous-market contract architecture.

7. Two market-structure facts that frame the bidding evidence

The two-tier CCGT bid curve and the day-ahead/intraday contrast. Every Big-4 CCGT fleet runs a two-tier day-ahead offer: a low *competing* tier near marginal cost and a high *parked* tier above the clearing-price ceiling. Withholding is firm-wide, not a Naturgy story. The same fleets bid a single competitive tier with negligible withholding in the intraday auction, after Fase I has cleared. A unit’s marginal cost does not change between two auctions hours apart. The parked tier is a placeholder rather than a Fase I bid: Fase I up-redispach is settled through a separate restriction offer submitted to REE.⁹ MTU15-DA grants CCGT a within-hour degree of freedom over its competing tier; the CCGT auction-cleared scale-up in § 4(iii) is the competing tier expanding into the four quarters of each peak hour.

⁹Procedure PO 3.2, settlement PO 14.4 ap. 20.1; see also CNMC (2025b,a). Earlier drafts of this memo conflated the two; the distinction justifies measuring bid shape and bid levels in the competing tier only, which is what the window-specific p_{90} bandwidth in § 1 achieves.

Market concentration migrated from wholesale to ancillary services. Wholesale-market HHI hovers around 1 100 over 2016–2024 (below the moderate-concentration threshold of 1 500) and falls mid-day when solar peaks (solar is less concentrated than the Big-3). Day-ahead restrictions upward HHI above 4 000; deviations upward above 6 000; tertiary upward around 5 000. Median of 4 firms with accepted upward restrictions bids. As renewables eroded wholesale-market power, technical and locational requirements concentrated the ancillary markets that absorb their variability.

References

- Kay H. Brodersen, Fabian Gallusser, Jim Koehler, Nicolas Remy, and Steven L. Scott. Inferring causal impact using Bayesian structural time-series models. *Annals of Applied Statistics*, 9(1): 247–274, 2015.
- Jiafeng Chen and Jonathan Roth. Logs with zeros? some problems and solutions. *Quarterly Journal of Economics*, 139(2):891–936, 2024. doi: 10.1093/qje/qjad054.
- CNMC. Resolución de 12 de junio de 2025: P.o. 14.4 *Liquidación de servicios de ajuste y pagos por capacidad*. BOE-A-2025-13076, 2025a.
- CNMC. Resolución de 6 de marzo de 2025: P.o. 3.2 *Resolución de restricciones técnicas*. BOE-A-2025-5342, 2025b.
- CNMC. Informe de supervisión del mercado peninsular mayorista al contado de electricidad. año 2023. Expediente IS/DE/013/24, <https://www.cnmc.es/expedientes/isde01324>, 2025c.
- Clément de Chaisemartin and Xavier D’Haultfoe uille. *Causal Inference with Differences-in-Differences: Credible Answers to Hard Questions*. Princeton University Press, 2026. Forthcoming.
- Joscha Märkle-Huß, Stefan Feuerriegel, and Dirk Neumann. Contract durations in the electricity market: Causal impact of 15-min trading on the EPEX. *Energy Economics*, 88:104760, 2020.